

INT8 vs F32 inference latency comparison

